

# HAWK : Performance Monitoring Hive

## Introduction



JunHo Cho

Data Analysis Platform Team



# Reasons related to Hive performance

- Inefficient hive query
- Not suitable configuration of hadoop, hive and system
- Bad distributed strategy of data
- System problems

# How to find out the reason?

- Hive
  - Hive plan
  - Hive log
- Hadoop
  - MapReduce monitoring
  - Hadoop job log
- System

# How to find out the reason ?

```
ABSTRACT SYNTAX TREE:
(TOK_QUERY (TOK_FROM (TOK_JOIN (TOK_TABREF emp e) (TOK_TABREF dept d) (= (. (TOK_TABLE_OR_COL e) deptno) (. (TOK_TABLE_OR_COL d) deptno)))) (TOK_INSERT
(TOK_DESTINATION (TOK_DIR TOK_TMP_FILE)) (TOK_SELECT (TOK_SELEXPR (. (TOK_TABLE_OR_COL e) empno)) (TOK_SELEXPR (. (TOK_TABLE_OR_COL e) ename)) (TOK_SELEXPR
(. (TOK_TABLE_OR_COL d) dname)))) (TOK_WHERE (= (. (TOK_TABLE_OR_COL d) deptno) 30))))
```

```
STAGE DEPENDENCIES:
Stage-1 is a root stage
Stage-0 is a root stage
```

```
STAGE PLANS:
```

```
Stage: Stage-1
```

## Map Reduce

Alias -> Map Operator Tree:

```
d
  TableScan
    alias: d
    Filter Operator
      predicate:
        expr: (deptno = 30)
        type: boolean
      Reduce Output Operator
        key expressions:
          expr: deptno
          type: int
        sort order: +
        Map-reduce partition columns:
          expr: deptno
          type: int
        tag: 1
        value expressions:
          expr: deptno
          type: int
          expr: dname
          type: string
```

```
e
  TableScan
    alias: e
    Reduce Output Operator
      key expressions:
        expr: deptno
        type: int
      sort order: +
      Map-reduce partition columns:
        expr: deptno
        type: int
      tag: 0
      value expressions:
        expr: empno
        type: int
        expr: ename
        type: string
```

Reduce Operator Tree:

```
Join Operator
  condition map:
    Inner Join 0 to 1
  condition expressions:
    0 {VALUE._col0} {VALUE._col1}
    1 {VALUE._col0} {VALUE._col1}
  handleSkewJoin: false
  outputColumnNames: _col0, _col1, _col4, _col5
  Filter Operator
    predicate:
      expr: (_col4 = 30)
      type: boolean
    Select Operator
      expressions:
        expr: _col0
        type: int
        expr: _col1
        type: string
        expr: _col5
        type: string
      outputColumnNames: _col0, _col1, _col2
    File Output Operator
      compressed: false
      GlobalTableId: 0
      table:
        input format: org.apache.hadoop.mapred.TextInputFormat
        output format: org.apache.hadoop.hive ql.io.HiveIgnoreKeyTextOutputFormat
```

```
Stage: Stage-0
```

## Fetch Operator

```
limit: -1
```

## # of MapReduce Job Data Source/Target Operators

- Hive
- Hive plan
- Hive log
- Hadoop
- MapReduce monitoring
- Hadoop job log
- System

# How to find out the reason ?

**essionStart** SESSION\_ID="nexr\_201111060248" TIME="1320509893783"  
QueryStart **QUERY\_STRING**="select ename from emp" **QUERY\_ID**="nexr\_20111106024848\_36387a27-924a-4d03-9614-e9ab5a86898a" TIME="1320509912081"  
**Counters plan**="{\"queryId\":\"nexr\_20111106024848\_36387a27-924a-4d03-9614-e9ab5a86898a\",\"queryType\":null,\"queryAttributes\":  
{\"queryString\":\"select ename from emp\"},\"queryCounters\":null,\"stageGraph\":  
{\"nodeType\":\"STAGE\",\"roots\":null,\"adjacencyList\":[]},\"stageList\":  
[\"stageId\":\"Stage-1\",\"stageType\":\"MAPRED\",\"stageAttributes\":null,\"stageCounters\":{}],\"taskList\":  
[\"taskId\":\"Stage-1\_MAP\",\"taskType\":\"MAP\",\"taskAttributes\":null,\"taskCounters\":null,\"operatorGraph\":  
{\"nodeType\":\"OPERATOR\",\"roots\":null,\"adjacencyList\":[\"node\":\"TS\_0\",\"children\":[\"SEL\_1\"],\"adjacencyType\":\"CONJUNCTIVE\"},  
\"node\":\"SEL\_1\",\"children\":[\"FS\_2\"],\"adjacencyType\":\"CONJUNCTIVE\"]}],\"operatorList\":  
[\"operatorId\":\"TS\_0\",\"operatorType\":\"TABLESCAN\",\"operatorAttributes\":  
{\"alias\":\"emp\"},\"operatorCounters\":{}],\"done\":\"false\",\"started\":\"false\"},  
\"operatorId\":\"SEL\_1\",\"operatorType\":\"SELECT\",\"operatorAttributes\":  
{\"expressions\":\"ename\"},\"operatorCounters\":{}],\"done\":\"false\",\"started\":\"false\"},  
\"operatorId\":\"FS\_2\",\"operatorType\":\"FILESINK\",\"operatorAttributes\":null,\"operatorCounters\":{}],\"done\":\"false\",\"started\":\"false\"}],\"done\":\"false\",\"started\":\"false\"}],\"done\":\"false\",\"started\":\"true\"}" TIME="1320509912095"  
TaskStart TASK\_NAME="org.apache.hadoop.hive ql.exec.MapRedTask" TASK\_ID="Stage-1"  
QUERY\_ID="nexr\_20111106024848\_36387a27-924a-4d03-9614-e9ab5a86898a" TIME="1320509912099"  
TaskProgress TASK\_HADOOP\_PROGRESS="2011-11-06 02:48:38,867 Stage-1 map = 0%, reduce = 0%" TASK\_NUM\_REDUCERS="0"  
TASK\_NAME="org.apache.hadoop.hive ql.exec.MapRedTask" TASK\_NUM\_MAPPERS="1"  
TASK\_COUNTERS="org.apache.hadoop.hive ql.exec.Operator\$ProgressCounter.CREATED\_FILES:0,Job  
Counters.SLOTS\_MILLIS\_MAPS:308,Job Counters.Launched map tasks:1,Job Counters.Data-local map tasks:1,Map-Reduce  
Framework.CPU\_MILLISECONDS:0" TASK\_ID="Stage-1" QUERY\_ID="nexr\_20111106024848\_36387a27-924a-4d03-9614-e9ab5a86898a" TASK\_HADOOP\_ID="job\_201111060121\_0001" TIME="1320509918868"  
...  
**TaskEnd** TASK\_RET\_CODE="0" TASK\_HADOOP\_PROGRESS="2011-11-06 02:48:41,986 Stage-1 map = 100%, reduce = 100%"  
TASK\_NUM\_REDUCERS="0" TASK\_NAME="org.apache.hadoop.hive ql.exec.MapRedTask" TASK\_NUM\_MAPPERS="1"  
TASK\_COUNTERS="org.apache.hadoop.hive ql.exec.Operator\$ProgressCounter.CREATED\_FILES:1,Job  
Counters.SLOTS\_MILLIS\_MAPS:8909,Job Counters.Total time spent by all reduces waiting after reserving slots (ms):0,Job  
Counters.Total time spent by all maps waiting after reserving slots (ms):0,Job Counters.Launched map tasks:1,Job Counters.Data-local  
map tasks:1,Job Counters.SLOTS\_MILLIS\_REDUCE:0,FileSystemCounters.HDFS\_BYTES\_READ:  
512,FileSystemCounters.HDFS\_BYTES\_WRITTEN:81235,FileSystemCounters.HDFS\_BYTES\_WRITTEN:  
37,org.apache.hadoop.hive ql.exec.MapOperator\$Counter.DESERIALIZE\_ERRORS:0,Map-Reduce Framework.Map input records:8,Map-  
Reduce Framework.Spilled Records:0,Map-Reduce Framework.CPU\_MILLISECONDS:0,Map-Reduce Framework.Map input bytes:  
0,Map-Reduce Framework.Map output records:0,Map-Reduce Framework.SPLIT\_RAW\_BYTES:173" TASK\_ID="Stage-1"  
QUERY\_ID="nexr\_20111106024848\_36387a27-924a-4d03-9614-e9ab5a86898a" TASK\_HADOOP\_ID="job\_201111060121\_0001"  
TIME="1320509926006"  
**QueryEnd** QUERY\_STRING="select ename from emp" QUERY\_ID="nexr\_20111106024848\_36387a27-924a-4d03-9614-e9ab5a86898a" QUERY\_RET\_CODE="0" QUERY\_NUM\_TASKS="1" TIME="1320509926006"  
Counters plan="{\"queryId\":\"nexr\_20111106024848\_36387a27-924a-4d03-9614-e9ab5a86898a\",\"queryType\":null,\"queryAttributes\":  
{\"queryString\":\"select ename from emp\"},\"queryCounters\":null,\"stageGraph\":  
{\"nodeType\":\"STAGE\",\"roots\":null,\"adjacencyList\":[]},\"stageList\":  
[\"stageId\":\"Stage-1\",\"stageType\":\"MAPRED\",\"stageAttributes\":null,\"stageCounters\":{}],  
\"CNTR\_NAME\_Stage-1\_REDUCE\_PROGRESS\":\"100\",\"CNTR\_NAME\_Stage-1\_MAP\_PROGRESS\":\"100\"},\"taskList\":  
[\"taskId\":\"Stage-1\_MAP\",\"taskType\":\"MAP\",\"taskAttributes\":null,\"taskCounters\":null,\"operatorGraph\":  
{\"nodeType\":\"OPERATOR\",\"roots\":null,\"adjacencyList\":[\"node\":\"TS\_0\",\"children\":[\"SEL\_1\"],\"adjacencyType\":\"CONJUNCTIVE\"},  
\"node\":\"SEL\_1\",\"children\":[\"FS\_2\"],\"adjacencyType\":\"CONJUNCTIVE\"]}],\"operatorList\":  
[\"operatorId\":\"TS\_0\",\"operatorType\":\"TABLESCAN\",\"operatorAttributes\":  
{\"alias\":\"emp\"},\"operatorCounters\":{}],\"done\":\"true\",\"started\":\"true\"},  
\"operatorId\":\"SEL\_1\",\"operatorType\":\"SELECT\",\"operatorAttributes\":  
{\"expressions\":\"ename\"},\"operatorCounters\":{}],\"done\":\"true\",\"started\":\"true\"},  
\"operatorId\":\"FS\_2\",\"operatorType\":\"FILESINK\",\"operatorAttributes\":null,\"operatorCounters\":{}],\"done\":\"true\",\"started\":\"true\"}],\"done\":\"true\",\"started\":\"true\"}],\"done\":\"true\",\"started\":\"true\"}" TIME="1320509926007"

- Hive
- Hive plan
- Hive log
- Hadoop
- MapReduce monitoring
- Hadoop job log
- System

# How to find out the reason ?

Hadoop job\_201110281201\_1043 on [tb301](#)

```
User: root
Job Name: SELECT unfold(tbl_x, 0.0, 0.0, 0.0, 0....tbl(Stage-1)
Job File: https://b301.servicelog/data/hadoop-root/mapred/staging/root/staging/job\_201110281201\_1043/job.xml
Submit Host: b301
Submit Host Address: 10.1.3.1
Job-ACLs: All users are allowed
Job Setup: Successful
Status: Succeeded
Started at: Wed Nov 02 11:31:26 KST 2011
Finished at: Wed Nov 02 11:32:32 KST 2011
Finished in: 1mins, 6sec
Job Cleanup: Successful
```

Kind	% Complete	Num Tasks	Pending	Running	Complete	Killed	Failed/Killed Task Attempts
map	100.00%	1	0	0	1	0	0 / 0
reduce	100.00%	1	0	0	1	0	0 / 0

	Counter	Map	Reduce	Total
org.apache.hadoop.hive.ql.exec.Operator\$ProgressCounter	CREATED_FILES	0	1	1
File Input Format Counters	Bytes Read	0	0	0
Job Counters	SLOTS_MILLIS_MAPS	0	0	48,671
	Launched reduce tasks	0	0	1
	Total time spent in waiting for the reducers after reserving slots (ms)	0	0	0
	Total time spent by all maps waiting after reserving slots (ms)	0	0	0
	Launched map tasks	0	0	1
	Data-local map tasks	0	0	1
	SLOTS_MILLIS_REDUCE	0	0	10,282
File Output Format Counters	Bytes Written	0	0	0
FileSystemCounters	FILE_BYTES_READ	0	671	671
	HDFS_BYTES_READ	1,796,025	0	1,796,025
	FILE_BYTES_WRITTEN	37,446	37,415	74,861
	HDFS_BYTES_WRITTEN	0	488	488
org.apache.hadoop.hive.ql.exec.MapOperator\$Counter	DESERIALIZE_ERRORS	0	0	0
	Map output materialized bytes	671	0	671
	Map input records	12,153	0	12,153
	Reduce shuffle bytes	0	671	671
	Spilled Records	3	3	6
	Map output bytes	656	0	656

- Hive
  - Hive plan
  - Hive log
- Hadoop
  - MapReduce monitoring
  - Hadoop job log
- System



# How to find out the reason ?

```

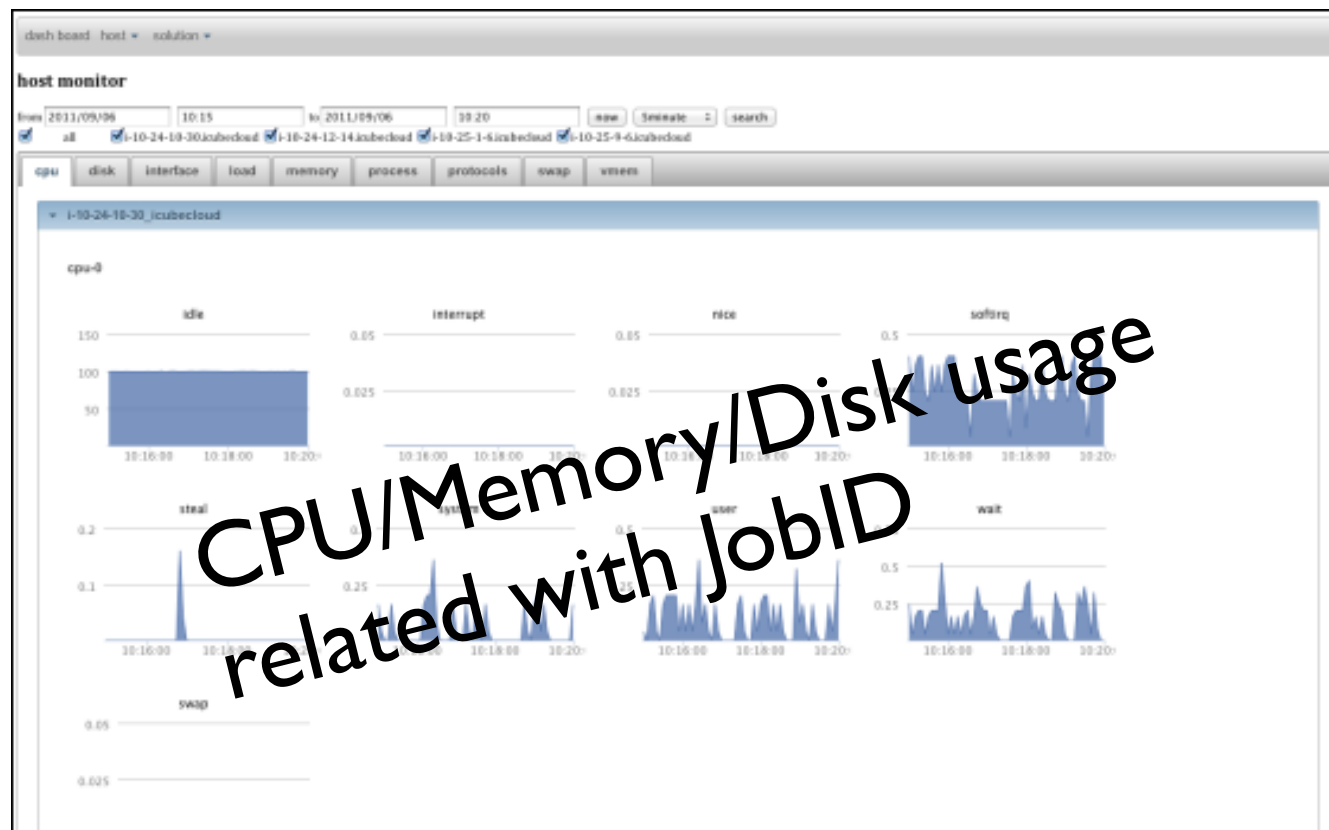
Meta VERSION="1".
Job JOBID="job_201111060121_0001" JOBNAME="select ename from emp(Stage-1)" USER="nextr" SUBMIT_TIME="1320509913049"
JOBCONF="hdfs://localhost:54310/tmp/hadoop-nextr/mapred/staging/nextr/\staging/job_201111060121_0001/job.xml" VIEW_JOB="*"
MODIFY_JOB="*" JOB_QUEUE="default".
Job JOBID="job_201111060121_0001" JOB_PRIORITY="NORMAL".
Job JOBID="job_201111060121_0001" LAUNCH_TIME="1320509913647" TOTAL_MAPS="1" TOTAL_REDUCES="0"
JOB_STATUS="PREP".

Task TASKID="task_201111060121_0001_m_000002" TASK_TYPE="SETUP" START_TIME="1320509913881" SPLITS="".
MapAttempt TASK_TYPE="SETUP" TASKID="task_201111060121_0001_m_000002"
TASK_ATTEMPT_ID="attempt_201111060121_0001_m_000002_0" START_TIME="1320509915561"
TRACKER_NAME="tracker_CentOS:CentOS/127.0.0.1:59707" HTTP_PORT="50060".
MapAttempt TASK_TYPE="SETUP" TASKID="task_201111060121_0001_m_000002"
TASK_ATTEMPT_ID="attempt_201111060121_0001_m_000002_0" TASK_STATUS="SUCCESS" FINISH_TIME="1320509918642"
HOSTNAME="/default-rack/CentOS" STATE_STRING="setup" COUNTERS="{(FileSystemCounters)(FileSystemCounters)
[(FILE_BYTES_WRITTEN)(FILE_BYTES_WRITTEN)(81235)]{(org.apache.hadoop.mapred.Task$Counter)(Map-Reduce
Framework)(SPILLED_RECORDS)(Spilled Records)(0)}"
Task TASKID="task_201111060121_0001_m_000002" TASK_TYPE="SETUP" TASK_STATUS="SUCCESS"
FINISH_TIME="1320509918715" COUNTERS="{(FileSystemCounters)(FileSystemCounters)[(FILE_BYTES_WRITTEN)
(FILE_BYTES_WRITTEN)(81235)]{(org.apache.hadoop.mapred.Task$Counter)(Map-Reduce Framework)(SPILLED_RECORDS)
(Spilled Records)(0)}"
Job JOBID="job_201111060121_0001" JOB_STATUS="RUNNING".
Task TASKID="task_201111060121_0001_m_000000" TASK_TYPE="MAP" START_TIME="1320509918718" SPLITS="/default-rack/
CentOS".
MapAttempt TASK_TYPE="MAP" TASKID="task_201111060121_0001_m_000000"
TASK_ATTEMPT_ID="attempt_201111060121_0001_m_000000_0" START_TIME="1320509918725"
TRACKER_NAME="tracker_CentOS:CentOS/127.0.0.1:59707" HTTP_PORT="50060".
MapAttempt TASK_TYPE="MAP" TASKID="task_201111060121_0001_m_000000"
TASK_ATTEMPT_ID="attempt_201111060121_0001_m_000000_0" TASK_STATUS="SUCCESS" FINISH_TIME="1320509921984"
HOSTNAME="/default-rack/CentOS" STATE_STRING="/rhive/data/emp1.txt:0-32" COUNTERS="{(org.apache.hadoop.hive ql.exec
\MapOperator$ProgressCounter)(org.apache.hadoop.hive ql.exec\Operator$ProgressCounter)[(CREATED_FILES)(CREATED_FILES)(1)]
{(FileSystemCounters)(FileSystemCounters)[(HDFS_BYTES_READ)(HDFS_BYTES_READ)(512)]{(org.apache.hadoop.hive ql.exec
\MapOperator$Counter)(org.apache.hadoop.hive ql.exec\MapOperator$Counter)[(DESERIALIZE_ERRORS)(DESERIALIZE_ERRORS)
(0)]{(org.apache.hadoop.mapred.Task$Counter)(Map-Reduce Framework)(MAP_INPUT_RECORDS)(Map input records)(8)]
[(SPILLED_RECORDS)(Spilled Records)(0)][(MAP_INPUT_BYTES)(Map input bytes)(0)][(MAP_OUTPUT_RECORDS)(Map output
records)(0)][(SPLIT_RAW_BYTES)(SPLIT_RAW_BYTES)(73)]"
Task TASKID="task_201111060121_0001_m_000000" TASK_TYPE="MAP" TASK_STATUS="SUCCESS" FINISH_TIME="1320509922187"
COUNTERS="{(org.apache.hadoop.hive ql.exec\Operator$ProgressCounter)(org.apache.hadoop.hive ql.exec\Operator
$ProgressCounter)[(CREATED_FILES)(CREATED_FILES)(1)]{(FileSystemCounters)(FileSystemCounters)[(HDFS_BYTES_READ)
(HDFS_BYTES_READ)(512)]{(FILE_BYTES_WRITTEN)(FILE_BYTES_WRITTEN)(81235)]{(HDFS_BYTES_WRITTEN)
(HDFS_BYTES_WRITTEN)(37)]{(org.apache.hadoop.hive ql.exec\MapOperator$Counter)(org.apache.hadoop.hive ql.exec
\MapOperator$Counter)[(DESERIALIZE_ERRORS)(DESERIALIZE_ERRORS)(0)]{(org.apache.hadoop.mapred.Task$Counter)(Map-
Reduce Framework)[(MAP_INPUT_RECORDS)(Map input records)(8)][(SPILLED_RECORDS)(Spilled Records)(0)]
[(MAP_INPUT_BYTES)(Map input bytes)(0)][(MAP_OUTPUT_RECORDS)(Map output records)(0)][(SPLIT_RAW_BYTES)
(SPLIT_RAW_BYTES)(173)]"
Task TASKID="task_201111060121_0001_m_000001" TASK_TYPE="CLEANUP" START_TIME="1320509922192" SPLITS="".
MapAttempt TASK_TYPE="CLEANUP" TASKID="task_201111060121_0001_m_000001"
TASK_ATTEMPT_ID="attempt_201111060121_0001_m_000001_0" START_TIME="1320509922195"
TRACKER_NAME="tracker_CentOS:CentOS/127.0.0.1:59707" HTTP_PORT="50060".
MapAttempt TASK_TYPE="CLEANUP" TASKID="task_201111060121_0001_m_000001"
TASK_ATTEMPT_ID="attempt_201111060121_0001_m_000001_0" TASK_STATUS="SUCCESS" FINISH_TIME="1320509924764"
HOSTNAME="/default-rack/CentOS" STATE_STRING="cleanup" COUNTERS="{(FileSystemCounters)(FileSystemCounters)
[(FILE_BYTES_WRITTEN)(FILE_BYTES_WRITTEN)(81235)]{(org.apache.hadoop.mapred.Task$Counter)(Map-Reduce Framework)
[(SPILLED_RECORDS)(Spilled Records)(0)]"
Task TASKID="task_201111060121_0001_m_000001" TASK_TYPE="CLEANUP" TASK_STATUS="SUCCESS"
FINISH_TIME="1320509925033" COUNTERS="{(FileSystemCounters)(FileSystemCounters)[(FILE_BYTES_WRITTEN)
(FILE_BYTES_WRITTEN)(81235)]{(org.apache.hadoop.mapred.Task$Counter)(Map-Reduce Framework)(SPILLED_RECORDS)
(Spilled Records)(0)]"

```

- Hive
  - Hive plan
  - Hive log
- Hadoop
  - MapReduce monitoring
  - Hadoop job log
- System

# How to find out the reason ?



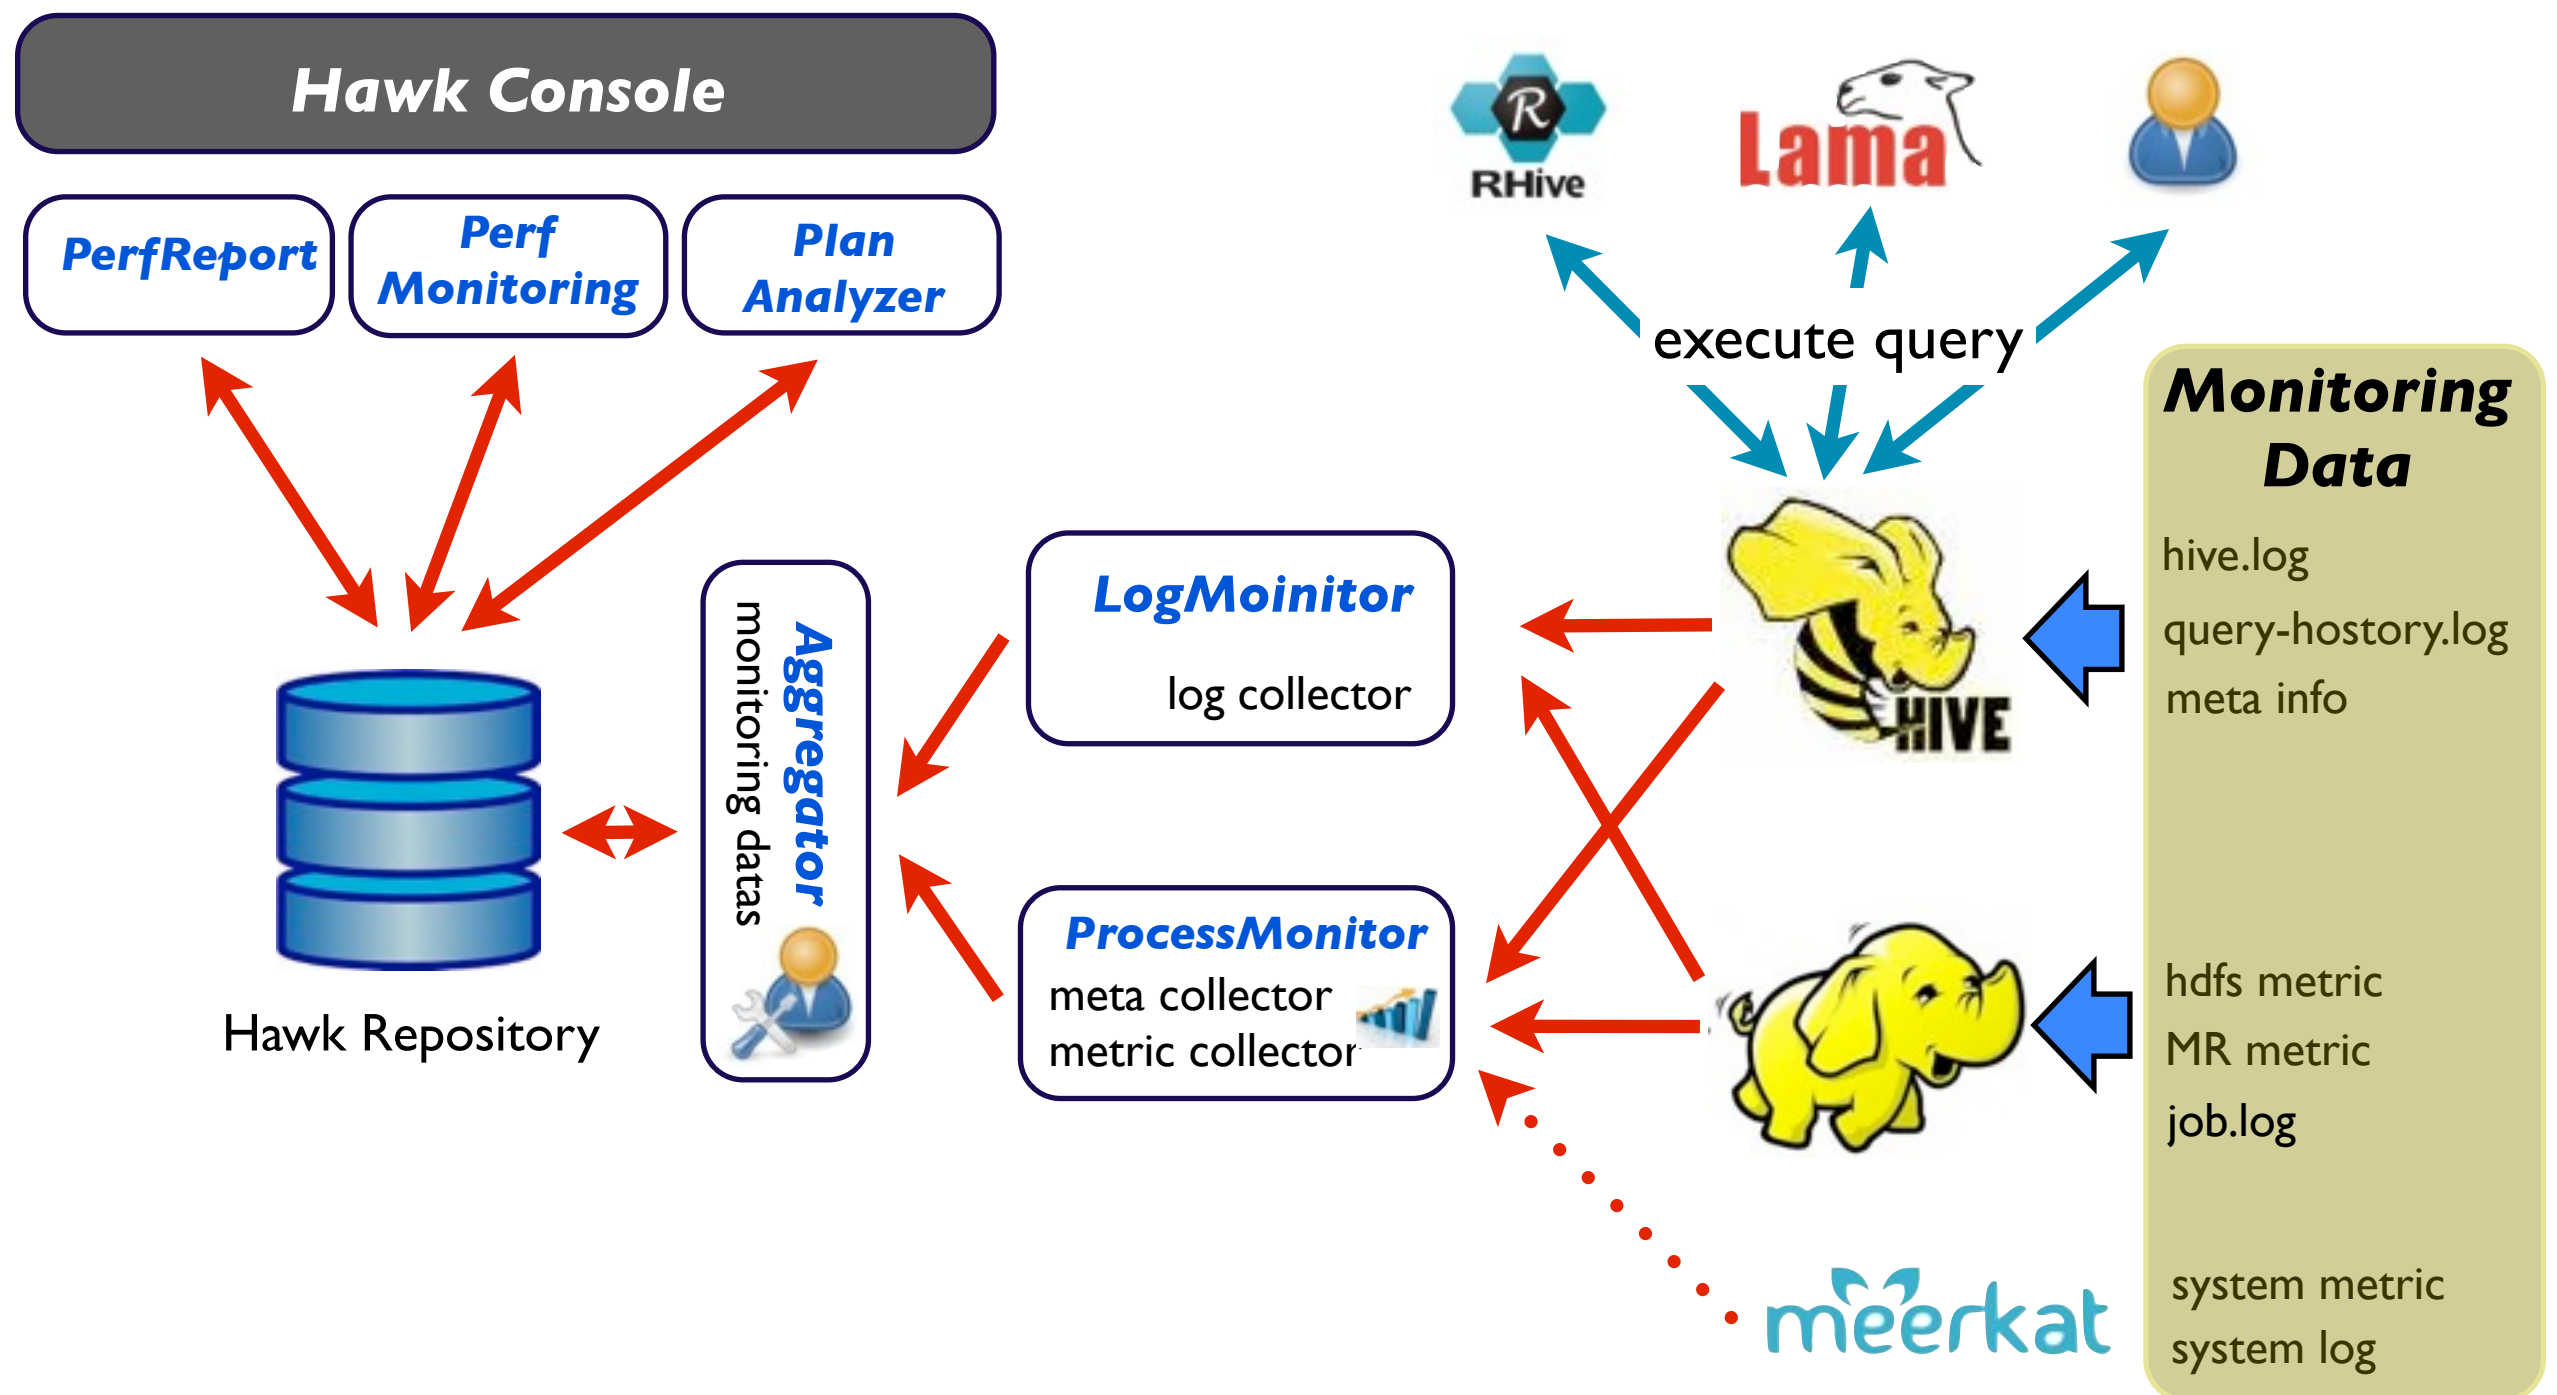
- Hive
  - Hive plan
  - Hive log
- Hadoop
  - MapReduce monitoring
  - Hadoop job log
- System



# Hawk Features

- Performance Data Collector
  - gathering hive / hadoop / sytem log
  - gathering hadoop / system metric
- Performance Report / Monitoring
  - Performance Report
    - hourly / daily / weekly / monthly report
  - Performance Analysis
    - Correlate related performance data (hive log, job log, mapreduce metric)
- Hive Plan Analyzer
  - Summarize hive system environment
  - Summarize hive query plan similar to Oracle plan
  - Correlate hive query plan to performance data
  - Query history management

# Architecture



# Hawk Plan Analyzer

hawk> plan nexr\_20111031065050\_548c8c6d-890c-46fa-a0a6-dfc6a0e51356

1

## Hive Configuration

```

134217728 = dfs.block.size
10737418240 = dfs.datanode.du.reserved
9 = dfs.datanode.handler.count
9 = dfs.namenode.handler.count
3 = dfs.replication
hive.hadoop.configurations = false
false = hive.auto.convert.join
false = hive.exec.compress.intermediate
false = hive.exec.compress.output
false = hive.exec.dynamic.partition
true = hive.exec.parallel
8 = hive.exec.parallel.thread.number
1000000000 = hive.exec.reducers.bytes.per.reducer
999 = hive.exec.reducers.max
org.apache.hadoop.hive.ql.io.CombineHiveInputF... = hive.input.format
100000 = hive.mapjoin.maxsize
25000000 = hive.mapjoin.smalltable.filesize
true = hive.mapred.reduce.tasks.speculative.execution
true = hive.optimize.cp
true = hive.optimize.groupby
true = hive.optimize.ppd
true = hive.optimize.pruner
false = hive.optimize.skewjoin
nexr_20111031065050_1ce9c1c9-6e59-4216-bed1-91... = hive.query.id
-1 = mapred.reduce.tasks
combiner_input_records / combiner_output_records
map_input_records / map_output_records
map_reduce_task_count
map_input_bytes / map_output_bytes
file_bytes_read / file_bytes_written
hdfs_bytes_read / hdfs_bytes_written
elapsed_time

```

2

## SQL

select deptno, count(1) cnt from emp E group by deptno

QUERY(nexr\_20111031065050\_70b58aaf-c160-4737-93b4-3d1f79eca4c5) - hash(850987f8417f1cedc5a9075bf442a476)

	ID	STEP	TABLE	MR JOBS	MR ROWS	MR BYTES	FILE	HDFS	TIME
hive stage	1	MAPRED (Stage-1)							00:00:16
mr job	2	MAP		1	11/5 (-/-)	-/76 (-)	-/84K	368/-	00:00:02
	3	TABLESCAN	emp (e)						-
	4	GROUPBY							-
mr task	5	REDUCE		1	5/-	-/(-)	92/83K	-/25	00:00:07
	6	GROUPBY	table name (alias)						-

3

## Predicate Information (Identified by operation ID) :

reduce\_shuffle\_bytes  
reduce\_input\_records / reduce\_output\_records

4

Last Executions : sqlhash=850987f8417f1cedc5a9075bf442a476

nexr\_20111031065050\_548c8c6d-890c-46fa-a0a6-dfc6a0e51356 | 2011-10-31 06:50:01.0 | 00:00:16 | select deptno, count(1) cnt from emp E ...

query\_id

start\_time

elapsed\_time

# UseCase

8 Nodes Hadoop Cluster  
# of Concurrent Reduce Task is 4

## Hive Query

```
CREATE TABLE TMP_VOICE
ROW FORMAT SERDE 'org.apache.hadoop.hive.serde2.columnar.ColumnarSerDe' STORED AS RCFILE
AS
SELECT yyyyymmdd, substr(yyyyymmdd,1,6) weekly, id, bits, call, rank
FROM (
  SELECT ((yyyyymmdd, weekly, id, bits, call, rank)
  FROM TMP_LOCAL_REAL_CUST k1
  UNION ALL
  SELECT yyyyymmdd, weekly, id, bits, call, rank
  FROM TMP_WORLD_REAL_CUST k2
) X
DISTRIBUTE BY HASH(X.yyyyymmdd,X.id)
SORT BY X.yyyyymmdd, X.id, X.call DESC , X.bits DESC
) t ) k2 WHERE rank <= 10;
```





# How to find out bottleneck using Hawk

## Hive Configuration

```

false = hive.auto.convert.join
false = hive.exec.compress.intermediate
false = hive.exec.compress.output
false = hive.exec.dynamic.partition
false = hive.exec.parallel
    
```

1000000000 = hive.exec.reducers.bytes.per.reducer

```

org.apache.hadoop.hive.ql.io.CombineHiveInputF... = hive.input.format
100000 = hive.mapjoin.maxsize
25000000 = hive.mapjoin.smalltable.filesize
true = hive.mapred.reduce.tasks.speculative.execution
true = hive.optimize.cp
true = hive.optimize.groupby
true = hive.optimize.ppd
true = hive.optimize.pruner
false = hive.optimize.skewjoin
-I = mapred.reduce.tasks
    
```

ID	STEP	TABLE	MR JOB	MR ROWS	MR BYTES	FILE	HDFS	TIME
1	MAPRED (Stage-1)							00:04:18
2	MAP		7	44.4M/44.4M(-/-)	973.0M/3.9G	3.9G/7.9G	1.5G/-	00:06:09
3	TABLESCAN	TMP_LOCAL_REAL_CUST (k1)						-
4	TABLESCAN	TMP_WORLD_REAL_CUST (k2)						-
5	REDUCE		2	44.4M/-	(3.3G)	4.0G/4.0G	-/1.6G	00:07:52
6	MOVE (Stage-0)							00:00:00
7	DDL (Stage-3)							00:00:01

Predicate Information (Identified by operation ID) :

```

5 - FILTER - (_col15 <= 10)
    
```

# How to find out bottleneck using Hawk

ID	STEP	TABLE	MR JOB	MR ROWS	MR BYTES	FILE	HDFS	TIME
1	MAPRED (Stage-1)							00:04:18
2	MAP		7	44.4M/44.4M(-/-)	973.0M/3.9G	3.9G/7.9G	1.5G/-	00:05:09
3	TABLESCAN	TMP_LOCAL_REAL_CUST (k1)						-
4	TABLESCAN	TMP_WORLD_REAL_CUST (k2)						-
5	REDUCE		2	44.4M/-	(3.3G)	4.0G/4.0G	-/1.6G	00:07:52
6	MOVE (Stage-0)							00:00:00
7	DDL (Stage-3)							00:00:01

Predicate Information (Identified by operation ID) :

5 - FILTER - (\_col15 <= 10)

- set 'hive.exec.reducer.bytes.per.reducer=10000000' (1GB -> 100MB)



ID	STEP	TABLE	MR JOB	MR ROWS	MR BYTES	FILE	HDFS	TIME
1	MAPRED (Stage-1)							00:01:56
2	MAP		7	44.4M/44.4M(-/-)	973.0M/3.9G	4.0G/7.9G	1.5G/-	00:06:19
3	TABLESCAN	TMP_LOCAL_REAL_CUST (k1)						-
4	TABLESCAN	TMP_WORLD_REAL_CUST (k2)						-
5	REDUCE		15	44.4M/-	(3.3G)	4.0G/4.0G	-/1.6G	00:22:45
6	MOVE (Stage-0)							00:00:00
7	DDL (Stage-3)							00:00:00



# Hawk Monitoring

```
hawk> history -day 10 -s elapsed_time -top 10 -desc
```

WorkName	Type	StartTime	Elapsed
nexr_20111031013333_347bd21a-809f-4b95-9335-2fb25ca48be9	HIVE_JOB	2011-10-31 01:33:26.0	10854 sec
nexr_20111031012525_8c4b97a7-fe2c-4c12-ab47-788c68f4438d	HIVE_JOB	2011-10-31 01:25:01.0	505 sec
nexr_20111103122727_613c6904-b95d-42ba-9dd0-8d10448132c1	HIVE_JOB	2011-11-03 12:27:44.0	494 sec
nexr_20111103114040_6d36c088-54c2-44f0-aabe-468dcc70cf34	HIVE_JOB	2011-11-03 11:40:45.0	470 sec
nexr_20111101220101_a7a0cec3-a4fe-4283-8def-e483528b4999	HIVE_JOB	2011-11-01 22:01:32.0	326 sec
nexr_20111102192727_78c44bb8-4a4b-4458-b97e-31f768b32fbf	HIVE_JOB	2011-11-02 19:27:06.0	318 sec
nexr_20111102193636_b4e32c87-989b-402b-8f00-68dfc3d452b0	HIVE_JOB	2011-11-02 19:36:24.0	284 sec
nexr_20111031083939_c551ccb7-5ce4-40cb-8669-046f23482dcb	HIVE_JOB	2011-10-31 08:39:58.0	244 sec
nexr_20111102182323_9f04514e-7da7-445d-9e7f-e8c0e2dbf5c8	HIVE_JOB	2011-11-02 18:23:08.0	236 sec
nexr_20111031074747_dc20db4e-62bb-4b5e-b52b-b85427402479	HIVE_JOB	2011-10-31 07:47:30.0	230 sec

```
hawk> metric -n nexr_20111103122727_613c6904-b95d-42ba-9dd0-8d10448132c1 -c MAP_INPUT_RECORDS,MAP_INPUT_BYTES,FILE_BYTES_READ
```

TimeStamp	MAP_INPUT_RECORDS	MAP_INPUT_BYTES	FILE_BYTES_READ
2011-11-03 12:27:50.0	0	0	0
2011-11-03 12:27:55.0	0	0	0
2011-11-03 12:28:00.0	781769	25622343	0
2011-11-03 12:28:05.0	30226426	891339946	0
2011-11-03 12:28:10.0	30226426	891339946	0
2011-11-03 12:28:15.0	47213583	1371490355	40632320
2011-11-03 12:28:20.0	131044562	3788847892	5418582484
2011-11-03 12:28:25.0	141485193	4093985561	5804067416
2011-11-03 12:28:30.0	141485193	4093985561	5804067416
2011-11-03 12:28:35.0	170018501	4942812887	7355840354
2011-11-03 12:28:40.0	170018501	4942812887	7355840354
2011-11-03 12:28:45.0	176952495	5138310773	8131961290



# Hawk Report

How many hive queries were launched last week ?

Day Of Week	HiveJob Count	MR MapTask Count	MR ReduceTask Count	Elapsed Time
Monday	2	664	297	1384 sec
Sunday	17	5	11	165 sec
Thursday	14	666	197	1152 sec
Tuesday	54	1993	552	3102 sec
Wednesday	144	3623	1009	17092 sec

# Hawk Report

Day	HiveJob Count	MR MapTask Count	MR ReduceTask Count	Elapsed Time
2011/10/26	45	389	872	46468 sec
2011/10/27	14	346	197	1152 sec
2011/10/30	17	5	11	165 sec
2011/10/31	72	1664	297	13980 sec
2011/11/01	54	1993	552	3102 sec
2011/11/02	99	26	37	624 sec

How many MapReduces ran last night ?

# Hawk Report

Day	HiveJob Count	MR MapTask Count	MR ReduceTask Count	Elapsed Time
2011/10/26	45	3522	272	26458 sec
2011/10/27	12	636	97	112 sec
2011/10/30	17	5	11	165 sec
2011/10/31	72	1664	297	13990 sec
2011/11/01	54	1993	552	3102 sec
2011/11/02	99	26	37	624 sec

How long did hive queries run 2 days ago?

# Hawk Report

Which hive queries ran longest yesterday ?

StartTime	WorkName	Type	Query	MR MapTask Count	MR ReduceTask Count	Elapsed Time
2011-11-01 09:30:29.0	nexr_20111101095404_f0096078-3820-49f7-aeb5-0ff42e862ef	HIVE_JOB	select E.net_ind_cd, count(1)	2	1	504 sec
2011-11-01 22:19:11.0	nexr_20111101221505_a80001099c-44f1-99f7-3d0176-f2f	HIVE_JOB	CREATE TABLE z_vip_scd_dd	2	2	254 sec
2011-11-01 22:19:45.0	nexr_20111101221909_c751e2-1aef-482a-2281-fdb8a267f18	HIVE_JOB	CREATE TABLE z_net_ind_cd_dd	1	1	247 sec
2011-11-01 22:11:11.0	nexr_20111101221111_da5771a4-c8c4-4f33-d9d-7aef1e8cfd16	HIVE_JOB	CREATE TABLE z_sqd_wcd_scd_dd	58	17	240 sec
2011-11-01 22:28:23.0	nexr_20111101222828_ad246f78-1170-429d-9807-de443c358699	HIVE_JOB	CREATE TABLE z_ipas_sbf_sqd_dd	678	182	220 sec
2011-11-01 22:06:58.0	nexr_20111101220606_05b94957-9017-41a0-a21d-708e437c2028	HIVE_JOB	CREATE TABLE z_sqd_call_ncn	130	37	194 sec
2011-11-01 10:30:07.0	nexr_20111101103030_e65e08a1-5cda-48ff-a34f-b11d04936b30	HIVE_JOB	select E.net_ind_cd, count(1)	3	2	170 sec
2011-11-01 22:25:58.0	nexr_20111101222525_fc192ca7-dab1-4165-94c7-16162574271d	HIVE_JOB	CREATE TABLE IF NOT EXISTS z_v	33	9	145 sec
2011-11-01 19:15:29.0	nexr_20111101191515_4f4b2e30-1517-4481-b24f-6d50d472f236	HIVE_JOB	SELECT COUNT(1) cnt FROM ( SEL	34	10	94 sec

# ***DEMO***

# Future Works

- GUI based hawk tool
- Integrate with meerkat for system monitoring
- Monitoring hive lock manager
- Performance Event alarm

# Data Analysis Platform Team

- JunHo Cho
- ByungMyon Chae
- Youngbae Ahn
- SeungWoo Ryu
- Seoeun Park
- Minwoo Kim
- Youngwoo Kim
- Seonghak Hong







[jun.cho@nexr.com](mailto:jun.cho@nexr.com)